

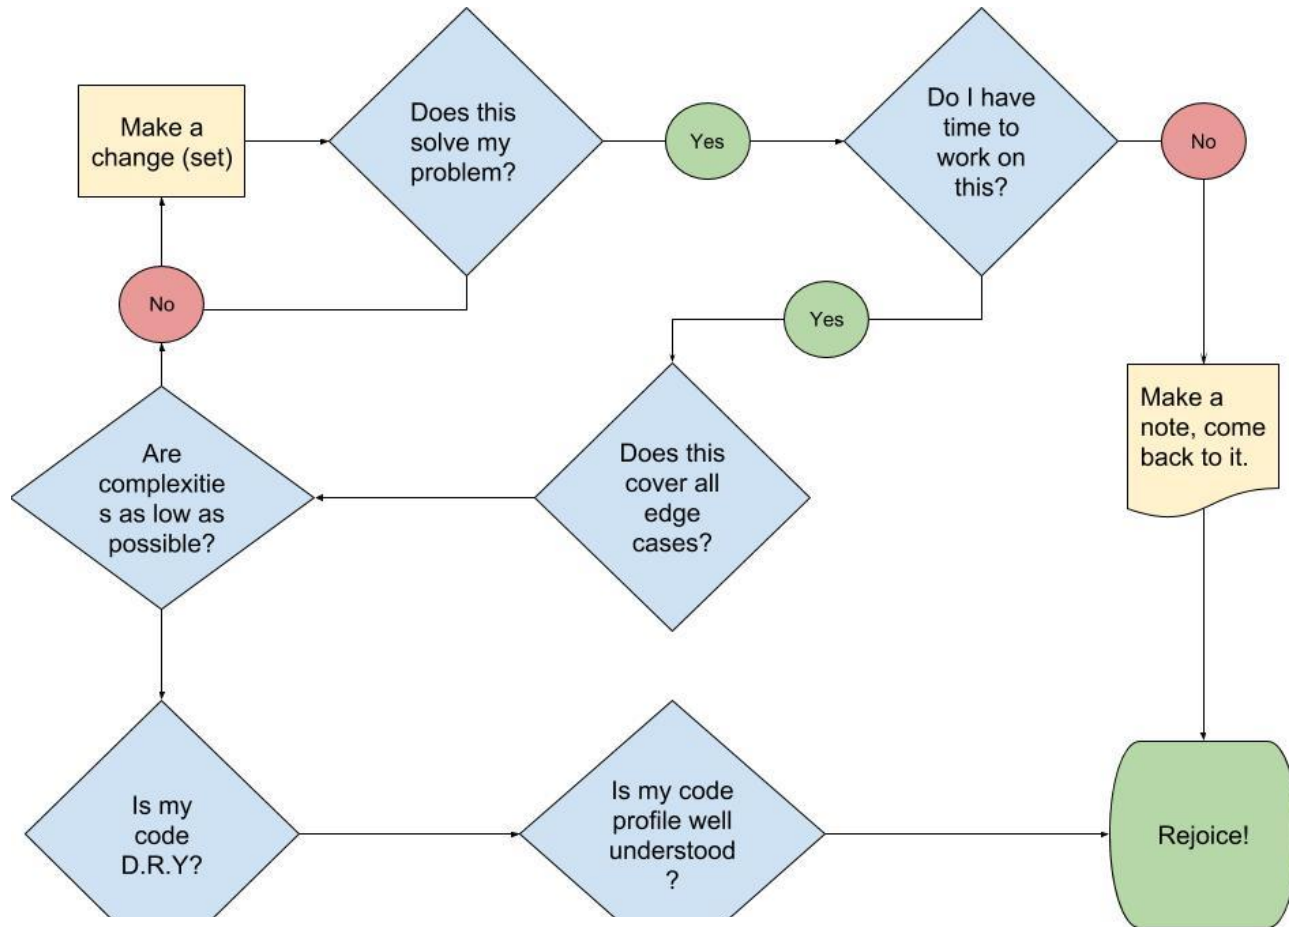


Code Profiling

(With a focus on Perl and Python)

Author: Joseph Wysk





Adapted from <https://medium.freecodecamp.org/time-is-complex-but-priceless-f0abd015063c>

Taking a step beyond time*

```
> which time  
/usr/bin/time  
> man time
```

“The time command runs the specified program command with the given arguments.

When command finishes, time writes a message to standard error giving timing statistics about this program run.

These statistics consist of

- (i) the elapsed real time between invocation and termination,
- (ii) the user CPU time (the sum of the tms_utime and tms_cutime values in a struct tms as returned by times(2))
- (iii) the system CPU time (the sum of the tms_stime and tms_cstime values in a struct tms as returned by times(2)).”



Perl

Perl: Out-of-the-box - Devel::DProf*

The produces a rather non-human friendly file, "tmon.out"

```
> perl -d:DProf example.pl
```

We can extract something human readable from tmon.out by simply calling

```
> dprofpp
```

Or you can cut to the chase with

```
> dprofpp -p example.pl
```

As you'd expect, you can learn about dprofpp via

```
> perldoc -F /usr/bin/dprofpp
```

<http://perldoc.perl.org/Devel/DProf.html>

*Although fun fact; this is deprecated and even Perldoc.org recommends using Devel::NTYProf

tmon.out:

```
#fOrTyTwo
$hz=100;
$XS_VERSION='DProf 20080331.00';
# All values are given in HZ
$over_untime=2; $over_stime=-1; $over_rtime=3;
$over_tests=10000;
$rrun_untime=10; $rrun_stime=0; $rrun_rtime=12;
$total_marks=2338
```

```
PART2
& 2 main BEGIN
+ 2
- 2
& 3 main BEGIN
+ 3
@ 0 0 1
& 4 strict bits
+ 4
- 4
& 5 strict import
+ 5
- 5
- 3
& 6 main BEGIN
+ 6
```

dprofpp:

Total Elapsed Time = 0.107659 Seconds
User+System Time = 0.117659 Seconds

Exclusive Times

%Time	ExclSec	Cumuls	#Calls	sec/call	Csec/c	Name
25.5	0.030	0.030	3	0.0100	0.0100	utf8::SWASHNEW
17.0	0.020	0.020	36	0.0006	0.0006	main::__ANON__
17.0	0.020	0.020	36	0.0006	0.0006	Safe::share_from
8.50	0.010	0.010	1	0.0100	0.0100	utf8::AUTOLOAD
8.50	0.010	0.010	15	0.0007	0.0006	Safe::BEGIN
8.50	0.010	0.029	36	0.0003	0.0008	Safe::rdo
0.00	--0.000	1	-	-	-	main::check_vol
0.00	--0.000	1	-	-	-	List::Util::bootstrap
0.00	--0.000	1	-	-	-	B::bootstrap
0.00	--0.000	1	-	-	-	subs::import
0.00	--0.000	1	-	-	-	Opcode::bootstrap
0.00	--0.000	1	-	-	-	Opcode::opset_to_ops
0.00	--0.000	1	-	-	-	UNIVERSAL::VERSION
0.00	--0.000	1	-	-	-	utf8::upgrade

Perl: w/ Devel::NYTProf

profile code and write database to ./nytprof.out

> perl -d:NYTProf some_perl.pl

convert database into a set of html files, e.g., ./nytprof/index.html

and open a web browser on the nytprof/index.html file

> nytprofhtml --open

or into comma separated files, e.g., ./nytprof/*.csv

> nytprofcsv

From: <https://metacpan.org/pod/Devel::NYTProf>

Perl: w/ Devel::NYTProf (cont'd.)

For a sufficient example we'll go on a field trip

<http://timbunce.github.io/devel-nytprof/sample-report/nytprof-20160319/index.html>

Perl: w/ Devel::NYTProf (cont'd.)

“The NYTProf profiler is written almost entirely in C and great care has been taken to ensure it's very efficient.” - <https://metacpan.org/pod/Devel::NYTProf>

In case this wasn't enough on it's own, there's a module extension for Apache

#Just add one line near the start of your httpd.conf file:

```
> PerlModule Devel::NYTProf::Apache
```

From: <https://metacpan.org/pod/Devel::NYTProf::Apache>

Taking a (perf)ect detour

```
> which perf?
```

“...perf is powerful: it can instrument CPU performance counters, tracepoints, kprobes, and uprobes (dynamic tracing). It is capable of lightweight profiling.

It is also included in the Linux kernel, under tools/perf, and is frequently updated and enhanced.

perf began as a tool for using the performance counters subsystem in Linux, and has had various enhancements to add tracing capabilities.” - https://perf.wiki.kernel.org/index.php/Main_Page

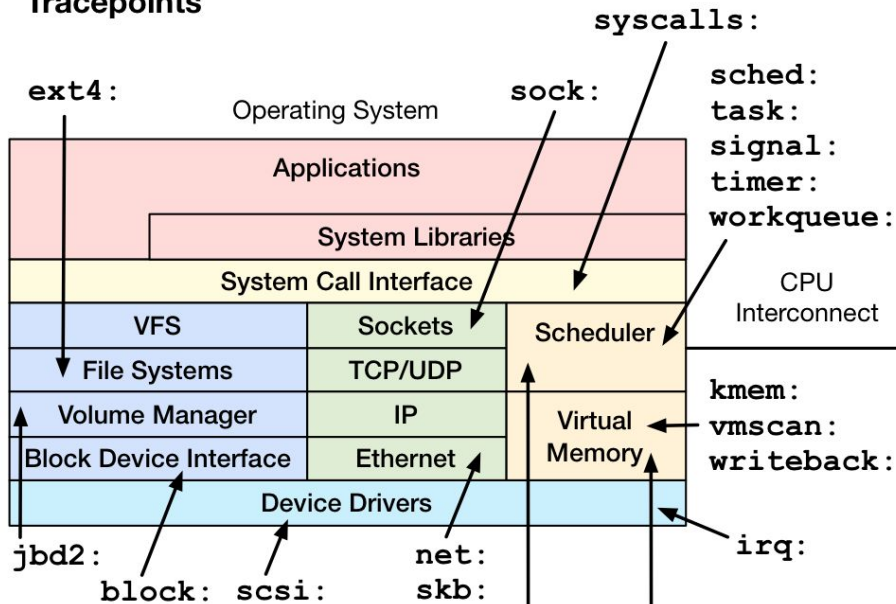
Linux perf_events Event Sources

Dynamic Tracing

uprobes

kprobes

Tracepoints



Software Events

cpu-clock
cs migrations

page-faults
minor-faults
major-faults

PMCs


cycles
instructions
branch-
L1-
LLC-*

CPU 1

Memory Bus


DRAM

mem-load
mem-store



perf record cmd and captured stack traces: ~8,000 lines

From: <http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html>



From: <http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html>

Perl: w/ Benchmark

Benchmark is a fantastic module for which I really should have an example.

Unfortunately, I don't at this time, so I highly encourage those interested to visit the metacpan page.

<https://metacpan.org/pod/Benchmark>



Python

Python: Out-of-the-box - cProfile

Basic usage

```
> import cProfile  
> cProfile.run('myFunction()', 'myFunction.profile')
```

We can extract something human readable via

```
> import pstats  
> stats = pstats.Stats('myFunction.profile')  
> stats.strip_dirs().sort_stats('time').print_stats()
```

You can also output a profile file for future use

```
> python -m cProfile -o example.prof example.py
```

Also see: <https://docs.python.org/2/library/profile.html>

Tue Sep 25 17:28:06 2018scicig_profiling.out

20096 function calls (20095 primitive calls) in 0.009 seconds

Ordered by: cumulative time

List reduced from 36 to 11 due to restriction <0.3>

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.000	0.000	0.009	0.009	scicig_profiling_example.py:3(<module>)
1	0.002	0.002	0.005	0.005	scicig_profiling_example.py:17(try_statement_test_case)
1	0.000	0.000	0.004	0.004	scicig_profiling_example.py:3(if_statement_test_case)
2/1	0.000	0.000	0.004	0.004	scicig_profiling_example.py:33(shuffle_array)
1	0.001	0.001	0.004	0.004	random.py:40(<module>)
10010	0.002	0.000	0.003	0.000	scicig_profiling_example.py:29(mk_array)
1	0.002	0.002	0.002	0.002	hashlib.py:56(<module>)
1	0.000	0.000	0.001	0.001	random.py:91(__init__)
1	0.000	0.000	0.001	0.001	random.py:100(seed)
10010	0.001	0.000	0.001	0.000	{method 'append' of 'list' objects}
1	0.000	0.000	0.000	0.000	{function seed at 0x7f6597c6b500}

Python: - PyCallGraph

This will generate an output *.png callgraph

```
> pycallgraph graphviz -o output_graph.png ./example.py
```

Perhaps use this if you want to bake something in or run on the fly

```
> from pycallgraph import PyCallGraph
```

```
> from pycallgraph.output import GraphvizOutput
```

From: <https://pycallgraph.readthedocs.io/en/master/>

```
from pycallgraph import PyCallGraph
from pycallgraph.output import GraphvizOutput

class Banana:
    def eat(self):
        pass

class Person:
    def __init__(self):
        self.no_bananas()

    def no_bananas(self):
        self.bananas = []

    def add_banana(self, banana):
        self.bananas.append(banana)

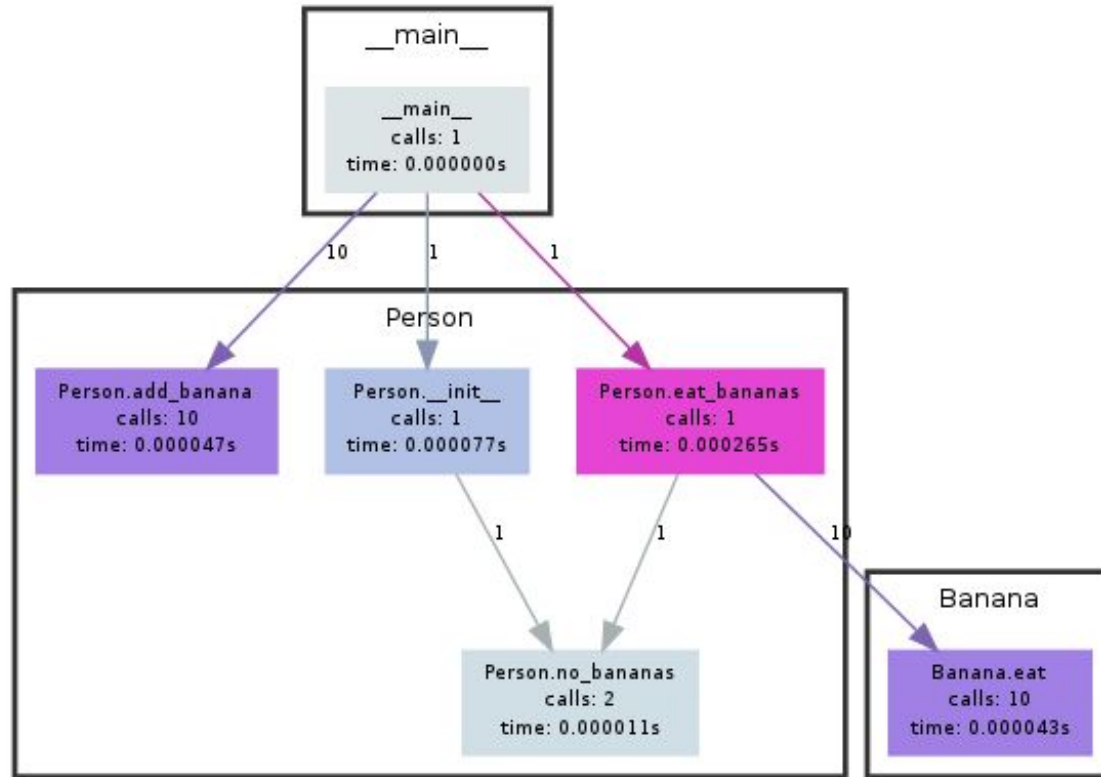
    def eat_bananas(self):
        [banana.eat() for banana in self.bananas]
        self.no_bananas()
)
```

```
def main():
    graphviz = GraphvizOutput()
    graphviz.output_file = 'basic.png'

    with PyCallGraph(output=graphviz):
        person = Person()
        for a in xrange(10):
            person.add_banana(Banana())
        person.eat_bananas()

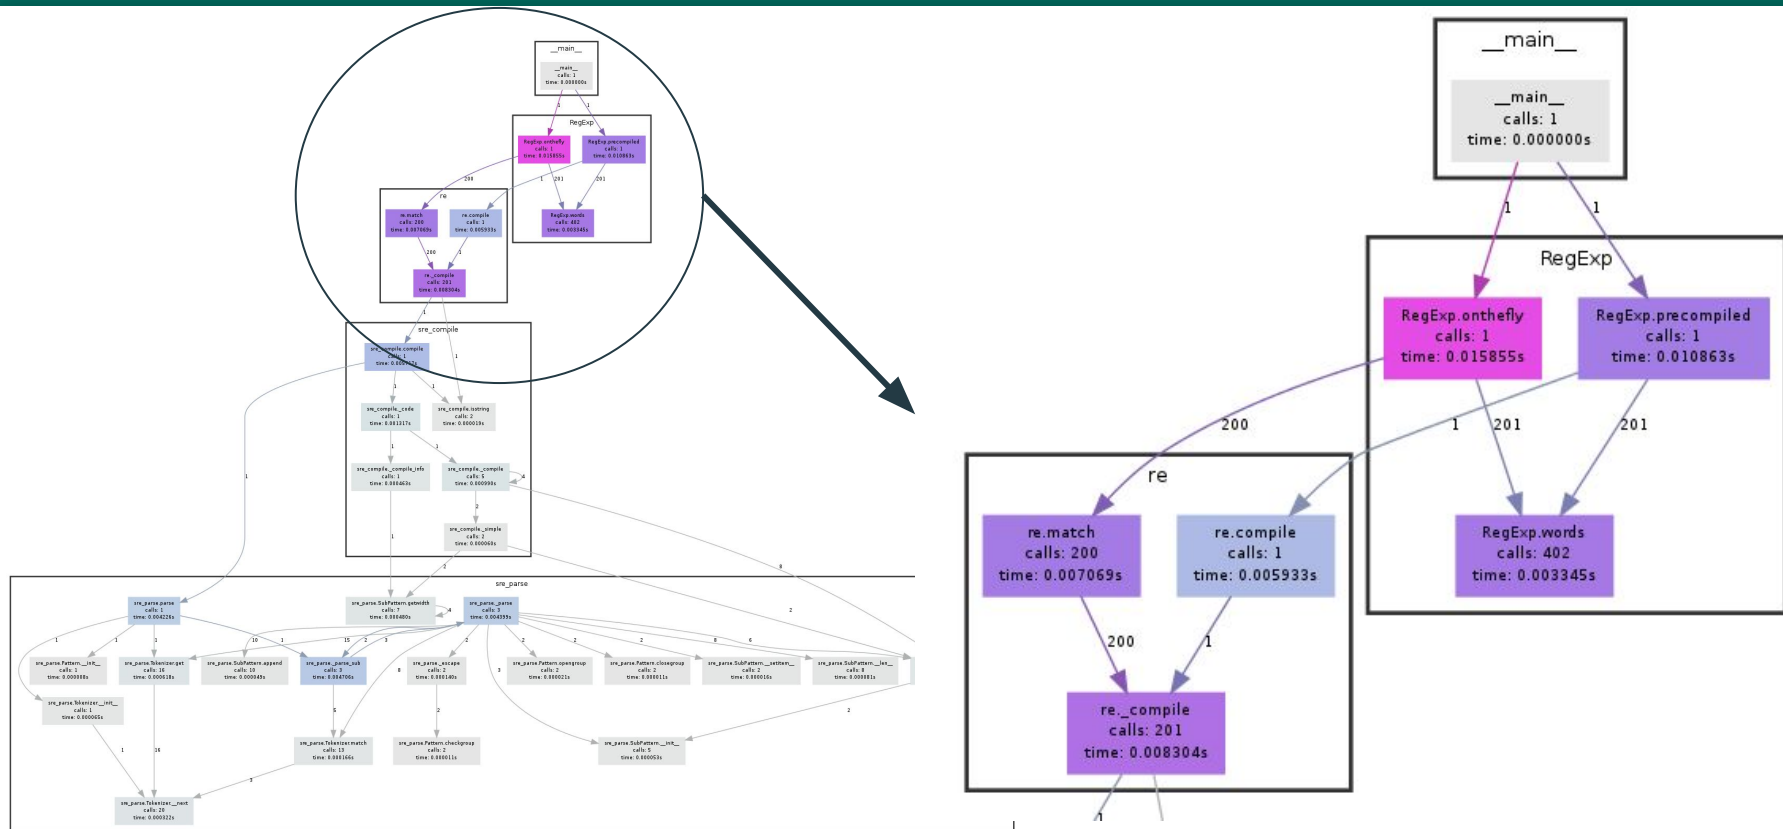
if __name__ == '__main__':
    main()
```

From: <https://pycallgraph.readthedocs.io/en/master/examples/basic.html>



Generated by Python Call Graph v1.0.0
<http://pycallgraph.slowchop.com>

<https://pycallgraph.readthedocs.io/en/master/examples/basic.html>



From: https://pycallgraph.readthedocs.io/en/master/examples/regexp_grouped.html

```
def if_statement_test_case():
    sample_array = []
    for x in xrange(0,10):
        if x != 0:
            value = x/x
            mk_array(sample_array, value)
        else:
            value = x*x
            mk_array(sample_array, value)

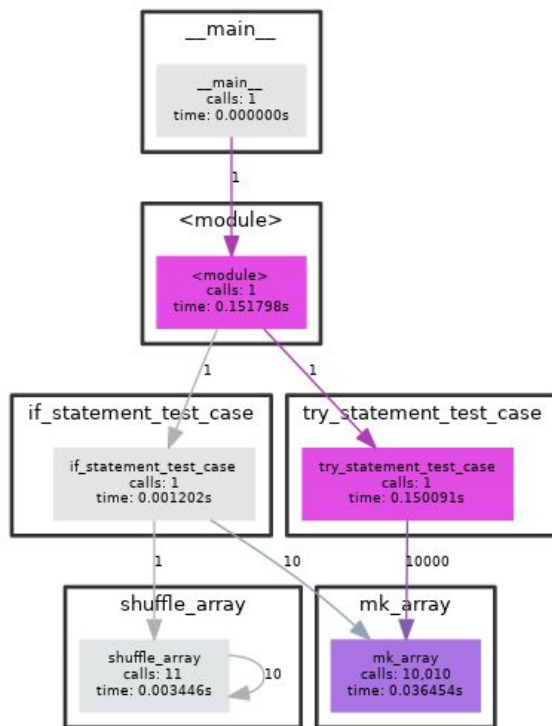
    shuffle_array(sample_array)
    return sample_array

def mk_array(arr,arr_input):
    arr.append(arr_input)
    return arr

def shuffle_array(arr):
    import random
    random.shuffle(arr)
    if arr[0] != 0:
        shuffle_array(arr)
    return arr
```

```
def try_statement_test_case():
    sample_array = []
    for x in xrange(0,10000):
        try:
            value = x/x
            mk_array(sample_array, value)
        except:
            value = x*x
            mk_array(sample_array, value)
    return sample_array

if __name__ == '__main__':
    if_statement_test_case()
    try_statement_test_case()
```



Generated by Python Call Graph v1.0.1
<http://pycallgraph.slowchop.com>

time_output:

real 0m0.276s

user 0m0.258s

sys 0m0.017s

> time pycallgraph graphviz -o ./myprofile.png ./scicig_profiling_example.py

Python: - timeit

- timeit is a very handy module for iterating over snippets of code to establish a consistent execution time

<https://docs.python.org/3/library/timeit.html>


```
from scicig_profiling_example import if_statement_test_case, try_statement_test_case

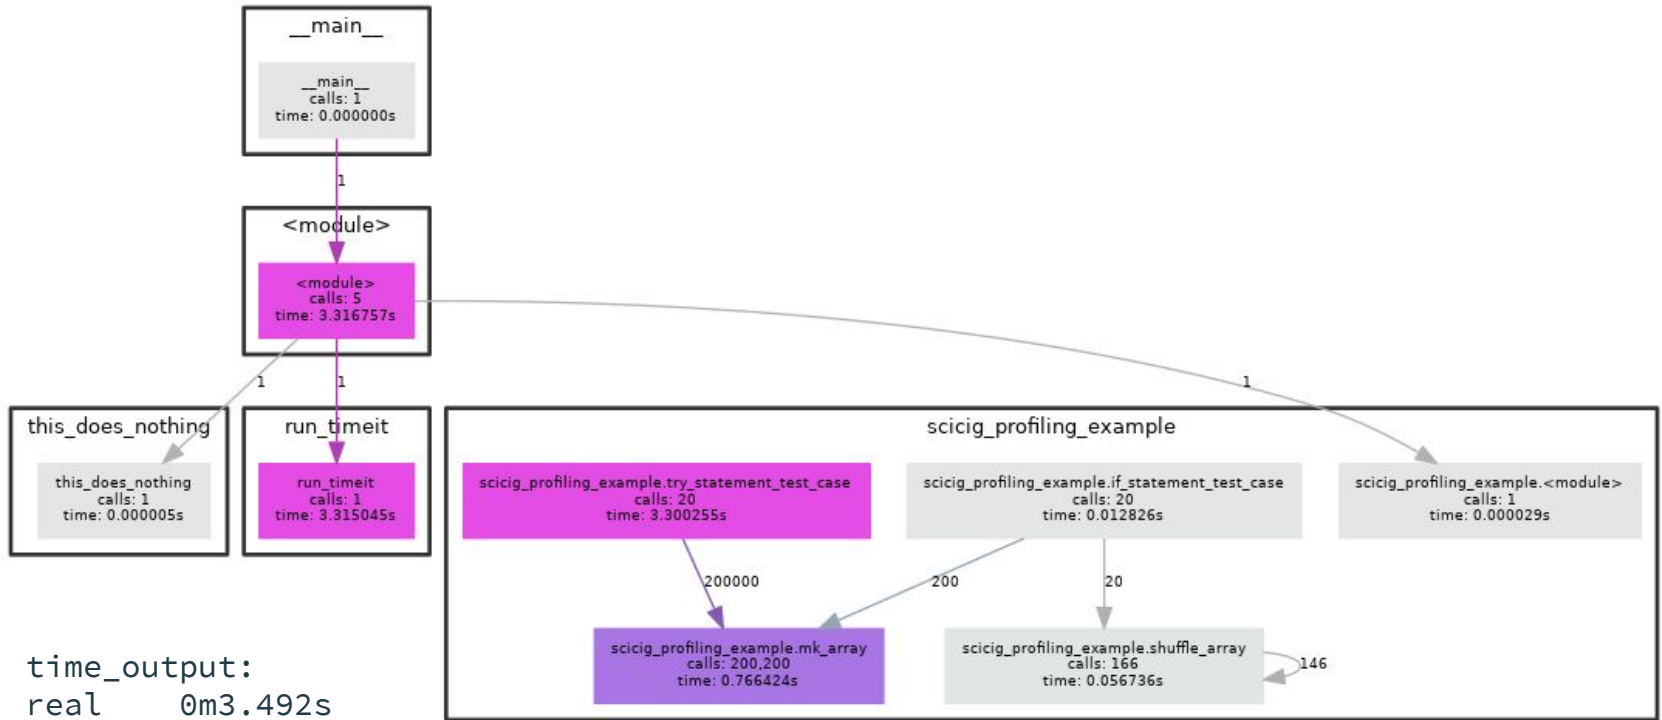
import timeit

def run_timeit():
    print min(timeit.repeat(if_statement_test_case, repeat=2, number=10))
    print min(timeit.repeat(try_statement_test_case, repeat=2, number=10))

def this_does_nothing():
    Return

this_does_nothing()

if __name__ == '__main__':
    run_timeit()
```



time_output:

real 0m3.492s

user 0m3.460s

sys 0m0.026s

Generated by Python Call Graph v1.0.1
<http://pycallgraph.slowchop.com>

```
> time pycallgraph graphviz -o ./timeit_profile.png ./timeit_test.py
```

```
timeit of if_statement_test_case = 0.00638699531555
```

```
timeit of try_statement_test_case = 1.64749479294
```

Python: - Anaconda Accelerate (Profiler)

This page on Anaconda Accelerate gives some good history and background to Numba/NumbaPro, which have/had an associated data profiler.

<https://www.anaconda.com/blog/developer-blog/open-sourcing-anaconda-accelerate/>

As far as I can tell, this data profiler is no longer support, but you can really do the same thing on your own with snakeviz.

<https://jiffyclub.github.io/snakeviz/>

Reset


Style: **Sunburst**

Depth: **20**

Cutoff: **1** / **1000**

Call Stack

0. scicig_profiling_example.py:3(<module>)



Search:

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.002181	0.002181	0.004935	0.004935	scicig_profiling_example.py:17(try_statement_test_case)
10010	0.00209	2.088e-07	0.002764	2.761e-07	scicig_profiling_example.py:29(mk_array)
1	0.00186	0.00186	0.001874	0.001874	hashlib.py:56(<module>)
1	0.000906	0.000906	0.003633	0.003633	random.py:40(<module>)

```
> python -m cProfile -o scicig_profiling.out ./scicig_profiling_example.py
> snakeviz -s -p 8888 timeit_test.out
```

Reset


Style: **Sunburst**

Depth: **20**

Cutoff: **1 / 1000**

Call Stack

```
2. scicig_profiling_example.py:33(shuffle_array)
1. scicig_profiling_example.py:31(statement_test_case)
0. scicig_profiling_example.py:3(<module>)
```



ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
1	0.002181	0.002181	0.004935	0.004935	scicig_profiling_example.py:17(try_statement_test_case)
10010	0.00209	2.088e-07	0.002764	2.761e-07	scicig_profiling_example.py:29(mk_array)
1	0.00186	0.00186	0.001874	0.001874	hashlib.py:56(<module>)
1	0.000906	0.000906	0.003633	0.003633	random.py:40(<module>)

Search:

Reset

Style:

Icicle

Depth:


5

Cutoff:

1

 /

1000



Search:


ncalls		tottime		percall		cumtime		percall		filename:lineno(function)
1		0.002181		0.002181		0.004935		0.004935		scicig_profiling_example.py:17(try_statement_test_case)
10010		0.00209		2.088e-07		0.002764		2.761e-07		scicig_profiling_example.py:29(mk_array)
1		0.00186		0.00186		0.001874		0.001874		hashlib.py:56(<module>)
1		0.000906		0.000906		0.003633		0.003633		random.py:40(<module>)

[Reset](#)

Style: **Sunburst**
Depth: **10**
Cutoff: **1** / **1000**

[Call Stack](#)

0. time_test_it.py:3(<module>)



Search:

ncalls		tottime		percall		cumtime		percall		filename:lineno(function)
20		0.03844		0.001922		0.0875		0.004375		scicig_profiling_example.py:17(try_statement_test_case)
200200		0.03731		1.864e-07		0.04914		2.454e-07		scicig_profiling_example.py:29(mk_array)
200204		0.01183		5.909e-08		0.01183		5.909e-08		~:0(<method 'append' of 'list' objects>)
1		0.001776		0.001776		0.001792		0.001792		hashlib.py:56(<module>)

```
> python -m cProfile -o timeit_test.out ./timeit_test.py
> snakeviz -s -p 8888 timeit_test.out
```



Great, so now what?

Python: - Numba

Numba leverages the LLVM compiler library to optimize machine code at runtime, claiming executions times comparable to C and FORTRAN.

Home: <https://numba.pydata.org/>

Quick start: <http://numba.pydata.org/numba-doc/latest/user/jit.html>

Examples: https://numba.pydata.org/numba-examples/examples/physics/lennard_jones/results.html

Python: - Numba

```
from numba import jit
@jit(nopython=True, parallel=True)
def jitsum(x):
    sum_val = 0
    for x in xrange(x):
        sum_val = sum_val + x
    return sum_val
%timeit jitsum(1000000)
1000000 loops, best of 3: 217 ns per loop
```

```
def mysum(x):
    sum_val = 0
    for x in xrange(x):
        sum_val = sum_val + x
    return sum_val
%timeit mysum(1000000)
10 loops, best of 3: 28.5 ms per loop
```

Python: - Dask

The emphasis for Dask is making parallel computation in python easy to leverage

Home: <http://dask.pydata.org/en/latest/>

Examples: <https://examples.dask.org/>

Tutorial: <https://github.com/dask/dask-tutorial>

Python: - Dask

```
import dask.array as da
```

```
x = da.random.random((10000, 10000),  
chunks=(1000, 1000))  
%timeit da.random.random((10000, 10000),  
chunks=(1000, 1000))  
100 loops, best of 3: 2.42 ms per loop
```

```
import numpy as np
```

```
x = np.random.rand(10000,10000)  
%timeit np.random.rand(10000,10000)  
1 loop, best of 3: 824 ms per loop
```

<https://towardsdatascience.com/how-i-learned-to-love-parallelized-applies-with-python-pandas-dask-and-numba-f06b0b367138>

Perl:

> use [Memoize](#);

“Memoizing' a function makes it faster by trading space for time. It does this by caching the return values of the function in a table. If you call the function again with the same arguments, memoize jumps in and gives you the value out of the table, instead of letting the function compute the value all over again.”

- <https://metacpan.org/pod/Memoize>

Advice

- Profiling and optimization tools can be easy to use, but are prone to being finicky
- Make this easy for yourself
- Automate and alias anything and everything
- Keep everything in perspective: Saving a second over a million iterations vs saving 15min every two weeks

“Indeed, one of my major complaints about the computer field is that **whereas Newton could say, ‘If I have seen a little farther than others, it is because I have stood on the shoulders of giants,’ I am forced to say, ‘Today we stand on each other’s feet.’”**

- Richard Wesley Hamming

https://en.wikiquote.org/wiki/Richard_Hamming